KRITIK SETH

New York, NY | (551)-344-6726 | sethkritik@gmail.com | linkedin.com/in/kritikseth | github.com/kritikseth

EDUCATION

New York University, Center for Data Science — **New York, USA**
Master of Science, Data Science (GenAI) | GPA: 3.8/4.0 — **May 2024**

- Relevant Courses: Machine Learning, Big Data, Computational Cognitive Modeling, Probability & Statistics, Natural Language Processing, Optimization & Computational Linear Algebra, Natural Language Understanding NLU, Advanced Python.

NMIMS University, MPSTME — **Mumbai, IN**
Bachelor of Technology, Data Science (Distinction) | GPA: 3.9/4.0 — **May 2022**

TECHNICAL SKILLS

- Programming & Tools: Python, R, SQL, PyTorch, TensorFlow, Scikit-Learn, Hugging Face, LangChain, Snowflake, Apache Spark, PySpark, Databricks, Airflow, Google Cloud Platform (GCP), Amazon Web Services (AWS), Tableau, Power BI, Excel
- ML & GenAI: Clustering, Regression, Classification, ANOVA, KNN, Decision Trees, Random Forest, K Means, Gradient Descent, Bagging, Boosting, Dimensionality Reduction, A/B Test, LLMs, RAG, Agentic Systems, Vector Databases, Transformer Models.

WORK EXPERIENCE

Bank of America – Software Engineer II (Customer Facing Contractor – ML) | Jersey City, New Jersey — **Sept 2024 – Present**

- Led GenAI / RAG workshops and demos (Gemini embeddings and Vector Search with Gemini LLM) and translated pilots into an agentic production workflow, ***cutting manual compliance review time by ~15% (≈2 hours/week saved per stakeholder)***, across frontline compliance teams.
- Partnered with Compliance and Trading to deliver POCs for a real-time anomaly detection pipeline (Vertex AI + BigQuery), demoing operational MLOps playbooks that drove adoption and secured production buy-in, helping ***reduce potential regulatory errors by ~5% on the monitored volume (1M+ trades/day)***, improving regulatory reporting accuracy.
- Deployed internal agentic workflow built with MCP server and Vertex AI to automate compliance logic and code reviews, proactively identifying data quality and rule gaps before trade submissions; reduced manual QA effort by 20% and improved cross-team collaboration between Compliance, Data, and Engineering.

NYU Stern School of Business – Data Science Project Lead | New York, New York — **May 2023 – June 2024**

- Led a team of 4 to develop the publicly operational Carbon Compass dashboard for NYC Local Law 97, integrating energy benchmarking and mortgage data in Tableau to ***reveal ~$450M in potential fines*** and promote citywide sustainability compliance.
- Cross-collaborated with policymakers, data engineers, and sustainability analysts to design scalable analytics workflows (Python, SQL, Tableau Prep, Airflow, AWS RDS/S3/Glue) that automated 11+ transformations, improved reporting turnaround by 40%.

Memorial Sloan Kettering Cancer Center – NLP Researcher | New York, New York — **Sept – Dec 2023**

- Fine-tuned LLMs and built an end-to-end ETL + inference pipeline (NER + classification) to automate detection and oncogenic classification of genes, outperforming LSTM baselines by ~5–6 pp while ***reducing inference time by 4×***.
- Integrated the system into OncoKB curation workflows; work was ***presented and published*** as a proffered poster at ***AACR 2025***.

Logitix – Data Science Intern | Boca Raton, Florida — **June – Dec 2023**

- ***Accomplished a 50% reduction in ticket pricing decision time (to 5 minutes)***, as measured by A/B testing, by training explainable ML models (SHAP), and deploying them via CI/CD as a real-time daily forecasting script integrated with the app.
- Trained an ensemble machine learning model (Random Forest and XGBoost) to predict ticket tiers with a 0.92 F1 score, securing lucrative partnerships with prestigious sports venues and directly ***generating $100K in revenue***.

Persistent Systems – Machine Learning Intern | Mumbai, India — **Jan – April 2022**

- Improved image segmentation efficiency by ***~15%, as measured by reducing labelling time from ~60 to ~50 minutes*** per batch, by applying hyperparameter tuning, cross-validation, and coordinating with technicians to refine model based on feedback.

PROJECTS

Instance-Level In-Context Unlearning in LLMs (Python, HuggingFace, LLaMA-3)

- Developed a targeted unlearning framework on the LLaMA-3 model with prefix and post-processing guardrail methods, achieving 87% forget accuracy while revealing "needle-in-a-haystack" scalability challenges in prompting-based unlearning.

Suspicious Clause Detection in T&C Documents (TensorFlow, HuggingFace, NLTK)

- Fine-tuned large language models (GPT) using TensorFlow and HuggingFace to build an NLP web application for detecting suspicious clauses in T&C documents, improving document review efficiency by significantly.

PUBLICATIONS

- OncoTagger: A novel LLM based tool for accelerated and accurate literature-based curation of cancer-related genes and variants. Research was published in *Cancer Research* and presented at the AACR Annual Meeting 2025. | **AACR 2025**
- Swachhdata: Open-source Python ETL library for text, image, and tabular data with ***100,000+ downloads*** | **PyPI, 2021–Present**